

From Diffuse Extension to Structured Scaffolding: Reframing the Extended Mind Thesis through Cognitive Architecture

Edervaldo José de Souza Melo

Independent Researcher

ORCID: 0009-0003-6835-135X

Email: edersouzamelo@gmail.com

December 21, 2025

Abstract

The Extended Mind thesis has played a central role in contemporary philosophy of cognition by challenging internalist conceptions of the mind and emphasizing the cognitive relevance of external artifacts. However, much of the debate has oscillated between internalist resistance and forms of externalism that risk conceptual diffuseness, treating cognitive extension as a loosely bounded and context-dependent phenomenon.

This paper proposes a reframing of the Extended Mind as *structured cognitive scaffolding*. On this view, cognitive extension is not merely a matter of functional coupling with external resources, but depends on the organization of stable, rule-governed, and functionally differentiated external structures that systematically participate in cognitive processes. Extension, therefore, is understood in architectural rather than incidental terms.

To clarify this proposal, the paper introduces *Nemosine* as a conceptual–operational case study. Nemosine is not presented as empirical evidence or as a universal model of cognition, but as an analytical device that renders visible how structured external scaffolds can support metacognition, regulation, and distributed cognitive control without collapsing into either internalism or naïve externalism.

By situating this analysis within the broader tradition of distributed cognition and Extended Mind research, the paper aims to refine the conceptual vocabulary of cognitive extension, highlighting the central role of structure, persistence, and architectural coherence in extended cognitive systems.

1 Introduction

The question of where cognition begins and ends has been a persistent concern in the philosophy of mind. Classical internalist approaches have traditionally located cognitive processes within the boundaries of the individual organism, typically identifying cognition with neural activity or, at most, with biologically constrained bodily processes. While such accounts offer clear ontological boundaries, they struggle to accommodate the pervasive role that external artifacts, representations, and environments play in everyday cognitive activity.

The Extended Mind thesis emerged as a direct challenge to this internalist picture. By arguing that cognitive processes can extend beyond the skin and skull when external resources play functionally equivalent roles to internal processes, proponents of the Extended Mind have significantly reshaped debates in cognitive science and philosophy. External notebooks, diagrams, tools, and digital artifacts are no longer seen merely as aids to cognition, but as potential constituents of cognitive systems themselves.

Despite its influence, the Extended Mind thesis has faced sustained resistance. Critics have questioned the criteria for cognitive inclusion, the stability of extended systems, and the risk of overextension, whereby almost any causal interaction with the environment could be labeled cognitive. In response, defenders of the thesis have often emphasized functional parity or causal integration, but these strategies have not fully resolved concerns about conceptual looseness and boundary indeterminacy.

This paper argues that a significant part of this impasse arises from treating cognitive extension primarily as a matter of *coupling* rather than *structure*. When extension is understood as an incidental or contextually defined relation between an agent and an external resource, the resulting accounts struggle to distinguish genuine cognitive integration from mere instrumental use. What is missing, the paper suggests, is an explicit focus on the architectural organization of extended cognitive systems.

To address this gap, the paper proposes a reframing of the Extended Mind in terms of *structured cognitive scaffolding*. On this view, external elements qualify as parts of an extended cognitive system not simply by being used or coupled, but by participating in stable, rule-governed, and functionally differentiated structures that support and constrain cognitive activity over time. Cognitive extension, therefore, is not an all-or-nothing phenomenon, but a matter of architectural integration.

Within this framework, the paper introduces *Nemosine* as a conceptual–operational case study. *Nemosine* is not presented as empirical evidence nor as a comprehensive model of cognition, but as an analytical device that illustrates how external scaffolds can be deliberately structured to support metacognition, regulation, and distributed cognitive control. Its role is illustrative rather than justificatory, serving to clarify the theoretical proposal advanced here.

The paper proceeds as follows. Section 2 revisits the core claims of the Extended Mind thesis and the main lines of criticism it has attracted. Section 3 develops the notion of structured cognitive scaffolding and articulates the criteria that distinguish scaffolds from mere tools. Section 4 situates this proposal within the broader tradition of distributed cognition and cognitive architecture. Section 5 examines *Nemosine* as a conceptual–operational case, highlighting how structured external systems can participate in cognition without inflating the scope of the mind. Section 6 discusses the implications of this reframing for Extended Mind theory, and Section 7 concludes by summarizing the contribution and outlining directions for further theoretical work.

2 The Extended Mind Thesis Revisited

The Extended Mind thesis was most prominently articulated by Clark and Chalmers (1998) as a challenge to the assumption that cognition is necessarily confined within the biological bound-

aries of the individual. Against internalist models that locate cognitive processes exclusively in neural or bodily mechanisms, the thesis proposes that external artifacts may, under certain conditions, constitute genuine parts of cognitive systems.

At the core of this proposal lies the principle of functional parity. According to this principle, if an external process plays the same functional role in a cognitive task as an internal process would, then there is no principled reason to deny it cognitive status. From this perspective, notebooks, diagrams, and other external resources can be seen not merely as aids to cognition, but as components of extended cognitive processes, provided that they are appropriately integrated into the agent's activity.

The appeal of the Extended Mind thesis lies in its capacity to account for the deep entanglement between agents and their environments in everyday cognition. However, this very permissiveness has also motivated sustained criticism. Opponents of the Extended Mind argue that functional parity alone is insufficient to demarcate genuinely cognitive processes from mere causal contributors. In particular, Adams and Aizawa contend that without principled constraints on cognitive inclusion, the boundaries of cognition risk becoming overly permissive, collapsing the distinction between cognition and background support Adams and Aizawa (2001, 2008).

Despite these strengths, the Extended Mind thesis has generated persistent controversy. One central line of criticism concerns the criteria for cognitive inclusion. Critics have argued that functional parity alone is insufficient to distinguish genuine cognitive components from mere causal contributors. If any resource that reliably influences behavior were granted cognitive status, the boundaries of cognition would become excessively permissive, threatening to collapse into triviality.

A related concern targets the stability and persistence of extended cognitive systems. Internal cognitive processes typically exhibit a degree of continuity, availability, and integration that is not always matched by external resources. Skeptics have therefore questioned whether transient or opportunistically used artifacts can plausibly be regarded as parts of cognitive systems, rather than as tools that merely support cognition from the outside.

In response to these objections, defenders of the Extended Mind have refined their position by emphasizing criteria such as reliable availability, automatic endorsement, and tight coupling between agent and artifact. These refinements aim to restrict cognitive extension to cases in which external resources are consistently and seamlessly integrated into cognitive routines, thereby excluding incidental or purely instrumental interactions.

While these responses address some of the most pressing objections, they do not fully resolve the underlying tension in the debate. The emphasis on coupling and integration, though necessary, leaves open questions about the internal organization of extended systems. What remains underdeveloped is a systematic account of how external components are structured, differentiated, and coordinated within extended cognitive architectures.

This limitation suggests that the central issue is not whether cognition can extend beyond the individual, but how such extension is organized. Without an explicit account of structure, Extended Mind theories risk oscillating between overly permissive externalism and restrictive internalist retreat. The next section develops an alternative framing that shifts the focus from

extension as mere coupling to extension as structured cognitive scaffolding.

3 From Extension to Structured Cognitive Scaffolding

Recent developments in the philosophy of mind have emphasized that cognitive extension should not be understood merely in terms of loose coupling between an agent and external resources, but rather in terms of stable forms of cognitive organization. In particular, Clark has argued that extended cognitive systems often rely on environmental structures that function as reliable components of cognitive routines, shaping problem-solving, planning, and self-regulation over time Clark (2008). This shift in emphasis motivates a reconceptualization of cognitive extension in architectural terms.

A key limitation of many Extended Mind accounts lies in their emphasis on *coupling* as the primary criterion for cognitive inclusion. While coupling captures the dynamic interaction between agents and external resources, it offers limited guidance for distinguishing between resources that merely assist cognition and those that participate in it as components of a cognitive system. Without further qualification, coupling risks collapsing into a broad notion of causal influence.

The concept of cognitive scaffolding provides a more precise lens through which to address this issue. Scaffolding, as used here, does not refer simply to any external aid or support, but to structured arrangements that actively shape, constrain, and stabilize cognitive activity. A scaffold is not merely something that an agent uses; it is something that reorganizes the space of possible cognitive actions by introducing durable constraints and affordances.

Crucially, cognitive scaffolds exhibit structural properties that distinguish them from transient tools. First, they are characterized by stability over time. Rather than being opportunistically recruited on a case-by-case basis, scaffolds persist as part of an agent's cognitive environment and are relied upon across multiple contexts. Second, scaffolds are rule-governed. Their operation is guided by explicit or implicit organizational principles that regulate how information is stored, accessed, and transformed. Third, scaffolds are functionally differentiated. Distinct components within the scaffold serve distinct cognitive roles, contributing to the coordination and regulation of cognitive processes.

These features mark a shift from viewing cognitive extension as an incidental relation to viewing it as an architectural phenomenon. On this view, external resources qualify as parts of extended cognitive systems not because they are frequently used, but because they are integrated into structured systems that support higher-order cognitive functions such as planning, monitoring, and self-regulation. Cognitive extension thus becomes a matter of system design rather than mere proximity or frequency of interaction.

This architectural emphasis resonates with broader developments in cognitive science that stress the role of environmental structure in shaping cognitive processes. In particular, work on distributed and situated cognition has shown that cognitive performance often depends on stable external arrangements that reduce internal computational demands and coordinate activity across time and contexts Hutchins (1995).

Importantly, this reframing does not reject the core insights of the Extended Mind thesis.

Instead, it refines them by introducing structural criteria that constrain when and how extension occurs. By focusing on scaffolding, the account avoids the permissiveness of treating any reliable coupling as cognitive, while also resisting a retreat to internalism. External structures can be genuinely cognitive, but only insofar as they participate in organized, persistent, and functionally articulated systems.

The notion of structured cognitive scaffolding therefore provides a principled middle ground in the debate over cognitive extension. It preserves the environmental openness emphasized by Extended Mind theorists while addressing longstanding concerns about boundary control and system integrity. The next section situates this proposal within the broader tradition of distributed cognition and cognitive architecture, further clarifying the relationship between scaffolding, structure, and extended cognitive systems.

4 Distributed Cognition and Cognitive Architecture

The proposal to reconceptualize the Extended Mind in terms of structured cognitive scaffolding finds strong affinities with earlier work in distributed cognition, most notably in the tradition articulated by Hutchins (1995). Rather than treating cognition as a property of isolated individuals, distributed cognition analyzes cognitive activity as emerging from systems composed of multiple interacting elements, including agents, artifacts, representations, and environmental structures.

A paradigmatic example of this perspective is found in the work of Hutchins (1995), who emphasizes that cognitive processes are often realized at the level of socio-technical systems rather than within individual minds. On this view, the relevant unit of analysis is not the isolated thinker, but the organized system through which information is propagated, transformed, and coordinated across internal and external resources. Importantly, such systems are characterized by stable representational formats and well-defined pathways of interaction.

What distinguishes distributed cognition from more permissive forms of externalism is its emphasis on structure. Cognitive systems are not merely collections of interacting components, but organized configurations in which roles are differentiated and processes are constrained by representational and procedural norms. Artifacts do not contribute to cognition simply by being present or causally influential, but by occupying specific positions within a coordinated architecture.

This architectural perspective aligns naturally with the notion of cognitive scaffolding developed in the previous section. Scaffolds can be understood as external components that are deliberately or historically stabilized within a cognitive system, shaping how information flows and how cognitive tasks are decomposed and managed. From this standpoint, scaffolding is not an auxiliary phenomenon, but a core mechanism by which distributed cognitive systems achieve reliability and control.

Cognitive architecture, in this context, refers to the organized arrangement of processes, representations, and control structures that govern cognitive activity. While traditional discussions of cognitive architecture have often focused on internal mechanisms, the distributed perspective highlights that architectural organization can span internal and external elements.

What matters is not the location of a component, but its functional role within the system as a whole.

By integrating insights from distributed cognition and cognitive architecture, the present framework clarifies how extended cognitive systems can maintain both openness and constraint. External scaffolds contribute to cognition not by blurring boundaries indiscriminately, but by participating in structured systems that exhibit persistence, coordination, and functional differentiation. This architectural emphasis provides the conceptual resources needed to move beyond debates over coupling and toward a more precise account of extended cognition.

The next section applies this framework to a conceptual–operational case study, examining how structured external scaffolding can be instantiated in practice without relying on empirical claims or general models of cognition.

5 Nemosine as a Conceptual–Operational Case

Nemosine is introduced here as a conceptual-operational case intended to clarify how structured external scaffolds can support higher-order cognitive organization without thereby trivializing the boundaries of cognition. The case is not offered as empirical evidence nor as a comprehensive model of cognition, but as an analytical device for making explicit the functional roles played by persistent, rule-governed external structures in metacognitive regulation, control, and coordination. For readers interested in additional background, supplementary materials are available in an open repository; however, such materials are not required for following the theoretical argument developed in this paper.

Nemosine is best understood as an externally structured cognitive environment designed to support higher-order cognitive activities such as planning, monitoring, and self-regulation. It does not aim to replace internal cognitive processes, nor does it function as an autonomous cognitive agent. Rather, it operates as a deliberately organized scaffold that shapes how cognitive tasks are externalized, decomposed, and coordinated over time.

From the perspective of structured cognitive scaffolding, the relevance of Nemosine lies in its architectural features rather than in any particular performance claims. The system is organized around persistent external structures that remain available across contexts, enabling continuity in cognitive activity. These structures are governed by explicit organizational principles that regulate how information is recorded, retrieved, and related, thereby introducing rule-governed constraints into the cognitive environment.

In addition, Nemosine exhibits functional differentiation among its components. Distinct external elements are associated with distinct cognitive roles, such as tracking commitments, organizing conceptual material, or supporting reflective monitoring. This differentiation allows the system to distribute cognitive control across specialized structures rather than concentrating it within a single locus. As a result, cognitive activity is guided not only by internal processes, but by the coordinated operation of external scaffolds.

Crucially, the role of Nemosine in this analysis is conceptual rather than justificatory. The system is not offered as evidence that cognition must be extended, nor as a benchmark for evaluating cognitive performance. Instead, it serves to illustrate how external resources can

be deliberately structured so as to participate in cognitive processes in a non-trivial way. Its significance lies in making visible the distinction between ad hoc tool use and organized cognitive scaffolding.

By examining Nemosine in this manner, the present account demonstrates how structured external systems can support extended cognition without inflating the scope of the mind or dissolving its boundaries entirely. The case thus helps to clarify what it means for external structures to be cognitively relevant within an architectural framework, setting the stage for a broader discussion of the implications of structured scaffolding for Extended Mind theory.

6 Implications for Extended Mind Theory

Reframing the Extended Mind in terms of structured cognitive scaffolding clarifies a long-standing tension in debates about the boundaries of cognition. One persistent concern has been whether extended accounts risk dissolving cognitive limits altogether, rendering cognition overly inclusive. In response to this worry, several authors have emphasized the importance of principled boundary conditions that distinguish genuine cognitive systems from loosely coupled supports Wilson (2002). The notion of structured scaffolding addresses this concern by grounding cognitive extension in stable, organized, and functionally integrated architectures rather than in mere causal interaction.

One immediate implication concerns the criteria for cognitive inclusion. Accounts that emphasize structured scaffolding provide a principled way to distinguish between genuine cognitive components and merely instrumental resources. External elements are not cognitively relevant simply because they are frequently used or causally efficacious, but because they occupy stable and functionally articulated roles within a broader cognitive architecture. This perspective constrains cognitive extension without reverting to internalist assumptions.

A second implication relates to the stability and persistence of extended cognitive systems. By emphasizing durable external structures and rule-governed organization, the scaffolding framework addresses concerns about the transient and opportunistic nature of many proposed cases of extension. Extended cognition is understood not as a fleeting interaction, but as a sustained mode of cognitive organization supported by persistent external scaffolds.

The architectural emphasis also clarifies the role of agency and control in extended systems. Rather than diffusing cognitive agency across an undefined environment, structured scaffolding highlights how control can be distributed across specialized components while remaining coordinated at the system level. This allows extended systems to exhibit coherence and regulation without postulating a single internal locus of control.

Importantly, these implications do not require abandoning the core insights of the Extended Mind thesis. The present framework preserves the thesis's emphasis on the openness of cognition to environmental resources, while refining the conditions under which such openness yields genuine cognitive extension. By introducing structure as a central criterion, the account avoids both excessive externalism and restrictive internalism.

Finally, the notion of structured cognitive scaffolding suggests a productive direction for future theoretical work. Rather than debating the metaphysical status of particular artifacts,

researchers can investigate the architectural properties of extended systems, examining how different forms of external organization support distinct cognitive functions. This shift promises a more tractable and theoretically grounded approach to understanding extended cognition.

7 Conclusion

This paper has argued for a reframing of the Extended Mind thesis in terms of structured cognitive scaffolding. While debates about cognitive extension have often focused on whether cognition can extend beyond the individual, the present account shifts attention to how such extension is organized once it is granted. By emphasizing structure rather than mere coupling, the paper offers a more precise way of distinguishing genuine cognitive integration from instrumental support.

The notion of structured cognitive scaffolding highlights the importance of stability, rule-governed organization, and functional differentiation in extended cognitive systems. External resources contribute to cognition not simply by being used or causally influential, but by occupying well-defined roles within persistent cognitive architectures. This perspective preserves the openness to environmental resources emphasized by the Extended Mind thesis while addressing longstanding concerns about boundary control and overextension.

By introducing Nemosine as a conceptual–operational case, the paper has sought to clarify these distinctions without relying on empirical claims or proposing a general model of cognition. The case serves to illustrate how external scaffolds can be deliberately structured to support higher-order cognitive functions while maintaining architectural coherence.

Taken together, these considerations suggest that the future of Extended Mind theory lies not in expanding the list of candidate artifacts, but in developing more refined accounts of the architectures through which cognition is extended. Focusing on structured scaffolding provides a promising framework for advancing theoretical discussions of extended cognition in a disciplined and productive direction.

References

- Adams, F. and Aizawa, K. (2001). The bounds of cognition. *Philosophical Psychology*, 14(1):43–64.
- Adams, F. and Aizawa, K. (2008). Why the mind is still in the head. *Philosophical Perspectives*, 22:1–21.
- Clark, A. (2008). *Supersizing the Mind: Embodiment, Action, and Cognitive Extension*. Oxford University Press, Oxford.
- Clark, A. and Chalmers, D. (1998). The extended mind. *Analysis*, 58(1):7–19.
- Hutchins, E. (1995). *Cognition in the Wild*. MIT Press, Cambridge, MA.
- Wilson, R. A. (2002). Bounded agency and the extended mind. *Synthese*, 129(1):91–123.